

Preliminary Processing of NGS Data

Mark Dunning

23 January 2018


Outline

This afternoon we will cover

- ▶ The *Fastq* format for sequencing reads
- ▶ Quality assessment of *fastq* files
- ▶ The *bam* format for representing *aligned* reads
- ▶ Practice using *Galaxy*


The sequencing process

- ▶ <https://www.youtube.com/embed/HMyCqWhwB8E> for a nice video
- ▶ Sequencing consists of a series of *cycles* (e.g. 100), at each cycle we try and incorporate different bases (A, T, C, G)
- ▶ Base that is successfully added will illuminate in a particular colour



Imaging

- ▶ Much of the sequencing time is spent taking images of the flowcell
- ▶ It is these images that are used to discover what fragments of DNA were sequenced
- ▶ This process is not perfect and can introduce *uncertainty*



Scale of data

Instrument	No. of Reads	Size
Ion Torrent PGM	5 million reads	1Gb
MiSeq	25 million reads	6GB
HiSeq rapid run	600 million	150GB (*)
HiSeq high-output	4 billion	1 TB

- Equivalent to **40** HD movies File sizes are for 100 bp reads, unzipped Number of reads from thermofisher.com and illumina.com

Fastq format

- ▶ A text file
- ▶ Can be *compressed* as a gz file
- ▶ Four lines per read

Sequence ID Sequenced Read

```
@SRR081708.237649/1  
GGGTTAGGGTTAGGGTTAGGGTTAGGGTTAGGGTTAGGGTTAGGGTTAGGGTTAGG  
+  
!>:<9>>>:=;>?<:@?>;==@@@>?=AAA<=A@?6>4B=<>>.@?<@;?#####
```

Blank Quality Score

Figure 3

Sequence names

Can contain the following

- ▶ Name of sequencer
- ▶ Flow cell *lane*
- ▶ Coordinates of the read on the flow cell
- ▶ whether this is a *paired* read and whether it is read 1, or 2

Quality scores

- ▶ Base-calling has some probability (p) that we make a mistake.
- ▶ The quality score expresses our *confidence* in a particular base-call; higher quality score, higher confidence
- ▶ One such score for each base of sequencing. i.e. 100 scores for 100 bases of sequencing
- ▶ These are of importance if we want to call SNVs etc.
 - ▶ need to be sure that differences detected from the reference genome are legitimate, and not caused by sequencing error

N?>:<9>>>:=;>>?<>:@?>;==@@@>?=AAA<>=A@?6>4B=<>>.@>?<@;?####

Quality Scores to probabilities

- ▶ look-up the ASCII code for each character
- ▶ subtract the offset to get the Q score
- ▶ convert to a probability using the formula:-

$$p = 10^{-Q/10}$$

Worked Example

for our particular example:

N?>:<9>>>:=;>>?<>:@?>;==@@@>?=AAA<>=A@?6>4B=<>>.@>?<@;?####

it works out as follows:-

	Character	Code	Minus.Offset..33..	Probability
1	N	78	45	0.00003
2	?	63	30	0.00100
3	>	62	29	0.00126
4	:	58	25	0.00316
5	<	60	27	0.00200
6	9	57	24	0.00398
7	>	62	29	0.00126
8	>	62	29	0.00126
9	>	62	29	0.00126
10	:	58	25	0.00316

...

Exercise

- ▶ Use the Galaxy tool *Text Manipulation* -> *Select last*
 - ▶ print the last 12 lines from the file
JoeBlogsBRCAPanel_R2.fastq
- ▶ How many reads are shown in the result?
- ▶ Look at the last read and write down the first five and last five ASCII characters
 - ▶ is the quality greater at the start, or the end of the read?
- ▶ Use the Galaxy tool *Text Manipulation* -> *Line/Word/Character count*
 - ▶ count how many lines are in the file
JoeBlogsBRCAPanel_R2.fastq in total
 - ▶ how many reads does this correspond to?

FastQC: Quality Assessment of fastqc

- ▶ FastQC from the Babraham Institute Bioinformatics Core has emerged as the standard tool for performing quality assessment on sequencing reads;
<https://www.bioinformatics.babraham.ac.uk/projects/fastqc/>
- ▶ The manual for fastqc is available online and is very comprehensive
- ▶ A “*traffic light*” system is used to draw your attention to sections of the report that require further investigation.
- ▶ fastqc will not actually *do* anything to your data. If you decide to trim or remove contamination for your samples, you will need to use another tool.
- ▶ it doesn't know what type of sequencing has been performed (WGS, exome, RNA-seq), which can affect interpretation of some of the plots

Example sections of a fastqc report

1. Basic Statistics



Basic Statistics

Measure	Value
Filename	sample.fastq
File type	Conventional base calls
Encoding	Illumina 1.5
Total Sequences	9053
Sequences flagged as poor quality	0
Sequence length	36
%GC	50

Figure 5

Example sections of a fastqc report

2. Per-base sequence quality





Figure 6

Example sections of a fastqc report

Ideally, the plot should look *something* like following:-



Per base sequence quality




Example sections of a fastqc report

3. Per-sequence quality scores



Per sequence quality scores



Exercise

- ▶ using the Galaxy Tool *NGS: QC and manipulation* -> *FastQC Read Quality reports*
 - ▶ generate a QC report for the file
JoeBlogsBRCAPanel_R2.fastq
- ▶ look at the basic statistics for the file
 - ▶ does the number of reads agree with your previous answer?
- ▶ is there any evidence for a decrease in quality as the read length increases?

The SAM / BAM format

- ▶ we don't really spend much time look at *fastq* files
- ▶ most of our time is spent with *aligned* reads
 - ▶ i.e. we have used some software to tell us whereabouts in the genome each read belongs to
 - ▶ we will have a go at this in the practical

The .sam file

- ▶ **Sequence Alignment/Map (sam)**
- ▶ The output from an aligner such as bwa
- ▶ Same format regardless of sequencing protocol (i.e. RNA-seq, ChIP-seq, DNA-seq etc)
- ▶ May contain un-mapped reads
- ▶ Potentially large size on disk; ~100s of Gb
- ▶ Official specification can be found online
<http://samtools.github.io/hts-specs/SAMv1.pdf>
- ▶ We normally work on a compressed version called a .bam file.
See later.

The .sam file

Comprises a *tab-delimited* section that describes the alignment of each sequence in detail.

1	10	2	3	4	5	6	7	8	9	11
SRR081708.237649	163	1	10003	6	1567M	=	10041	105		
GACCCTGACCCTAACCCCTGACCCTGACCCTAACCCCTGACCCTAACCCCTAACCCCTAA S=<====<=>=<?=?>==@??;?>@@@=?@?@??@??@?>?@<@>@'@=?=?										
=<=>?>?>Q ZA:Z:<@;0;0;;308;68M;68><@;0;0;;27;;>MD:Z:5A11A5A11A5A11A13 RG:Z:SRR081708 NM:1:6 OQ:Z:GEGFFEGGGDGGGGDGA?										
DCDD:GGGDDGDCFGDFFCCCBEBFDABDD-D:EEEE=D=DDDDC:										

Figure 9

- ▶ 1:- Sequence ID
- ▶ 2:- Sequence quality expressed as a bitwise *flag*
- ▶ 3:- Chromosome that the read aligned to

Fun with flags!

The “*flags*” in the sam file can represent useful QC information

- ▶ Read is unmapped
- ▶ Read is paired / unpaired
- ▶ Read failed QC
- ▶ Read is a PCR duplicate (see later)

[illegible]

Figure 10

Derivation

	ReadHasProperty	Binary	MultiplyBy
Paired?	TRUE	1	1
Properly Paired?	TRUE	1	2
Unmapped?	FALSE	0	4
Unmapped Mate	FALSE	0	8
On Minus Strand?	FALSE	0	16
Mate on Minus Strand?	TRUE	1	32
Is First Read?	FALSE	0	64
Is Second Read?	TRUE	1	128
Is Secondary Alignment?	FALSE	0	256
Is Not Passing QC?	FALSE	0	512
Is Duplicate Read?	FALSE	0	1024

$$1 \times 1 + 1 \times 2 + 0 \times 4 + 0 \times 8 + 0 \times 16 + 1 \times 32 + 0 \times 64 + 1 \times 128 + 0 \times 256 + 0 \times 512 + 0 \times 1024 = 163$$

<https://broadinstitute.github.io/picard/explain-flags.html>

The .sam file

1 **2** **3** **4** **5** **6** **7** **8** **9** **11**

```
SRR081708.237649      163    1     10003   6       1567M   =          10041   105  
GACCTCGACCCTAACCTGACCTTGACCTGAAGCCTGACCTGACCTTAAACCTGTAA S<=====><?=?==@???:?@@@??@?????@?@?>?@&-@-'@=?=?  
=<=>?~?=Q           ZA:Z:<6;0;0;308;68M;68>;0;0;27;*>MD:Z:5A11A5A11A5A11A13   RG:Z:SRR081708 NM:i:6 OQ:Z:GEGFFFGGGDGDGGDDGA?  
DCDD:GGDGDCGFDFDFFFCCBBFDBADD-D:EED=D=DDDDC|]
```

Figure 11

- ▶ 4:- Start Position
- ▶ 5:- Mapping Quality; Confidence that an alignment is correct
- ▶ 6:- CIGAR; Describes positions of matches, insertions, deletions w.r.t reference

Have a CIGAR!

[illegible]

Figure 12

The **CIGAR** (**C**ompact **I**diosyncratic **G**apped **A**lignment **R**eport) string is a way of encoding the match between a given sequence and the position it has been assigned in the genome. It is comprised by a series of letters and numbers to indicate how many consecutive bases have that mapping.

- ▶ 68M
 - ▶ 68 bases matching the reference
- ▶ 1S67M
 - ▶ 1 soft-clipped read followed by 67 matches
- ▶ 15M87N70M90N16M
 - ▶ 15 matches following by 87 bases skipped followed by 70 matches etc.

The .sam file

[illegible]

Figure 13

7, 8, 9:- Alignment information for the paired read

The .sam file

1 **2** **3** **4** **5** **6** **7** **8** **9** **11**

SRR081708.237649 10 163 1 10003 6 1567M = 10041 105
GACCTTGACCTTAACCTGACCTTACCCTGACCTGACCTTAACCTGACCTTAACCTAA S=====<>?=?==@??;?@@@??@???????@?@?@?@?@'@=?=?
<====>=7=Q ZA:Z:<6;0;;308;68M;68><0;0;0;27; >MD:Z:5A11A5A11A5A11A13 RG:Z:SRR081708 NM:1:6 OQ:Z:GEGFFEGGGDGDGGDGGA?
DCDD:GGDGDGCGFDDFFFFCCBEFBADBD-D:EEEE=D=DDDDC;

Figure 14

- ▶ 10:-Sequence
- ▶ 11:- Base Qualities

This is the same as the `fastq` file; so if you have aligned data you can always go back and re-align